

Topic Models & Recurrent Neural Networks

Kuan-Yu Chen (陳冠宇)

2018/04/12 @ TR-409, NTUST

Semantic?

- In the context of NLP, long-span information is important!
 - Especially for language modeling

一 枝 美 麗 的 玫 瑰 在 花 盆

一 隻 可 愛 的 小 貓 在 花 園

一 支 高 貴 的 鋼 筆 在 拍 賣

- N-gram models are not an efficient strategy to capture the long-span information
 - From literal matching to semantic mapping

Probabilistic Latent Semantic Analysis

- **Probabilistic Latent Semantic Analysis** also called
 - Probabilistic Latent Semantic Indexing (PLSI)
 - Aspect Model
- PLSA is a probabilistic counterpart of LSA
 - $P(d_j)$: the probability of selecting document d_j
 - $P(w_i|T_k)$: the probability of word w_i condition on a latent factor/topic T_k
 - **Aspect!**
 - $P(T_k|d_j)$: the probability of a latent factor/topic T_k generated by document d_j

PLSA – 1

- The PLSA model is a latent variable model for co-occurrence data (i.e., each pair of word w_i and document d_j) which associates an unobserved class variable (i.e., latent factor T_k)

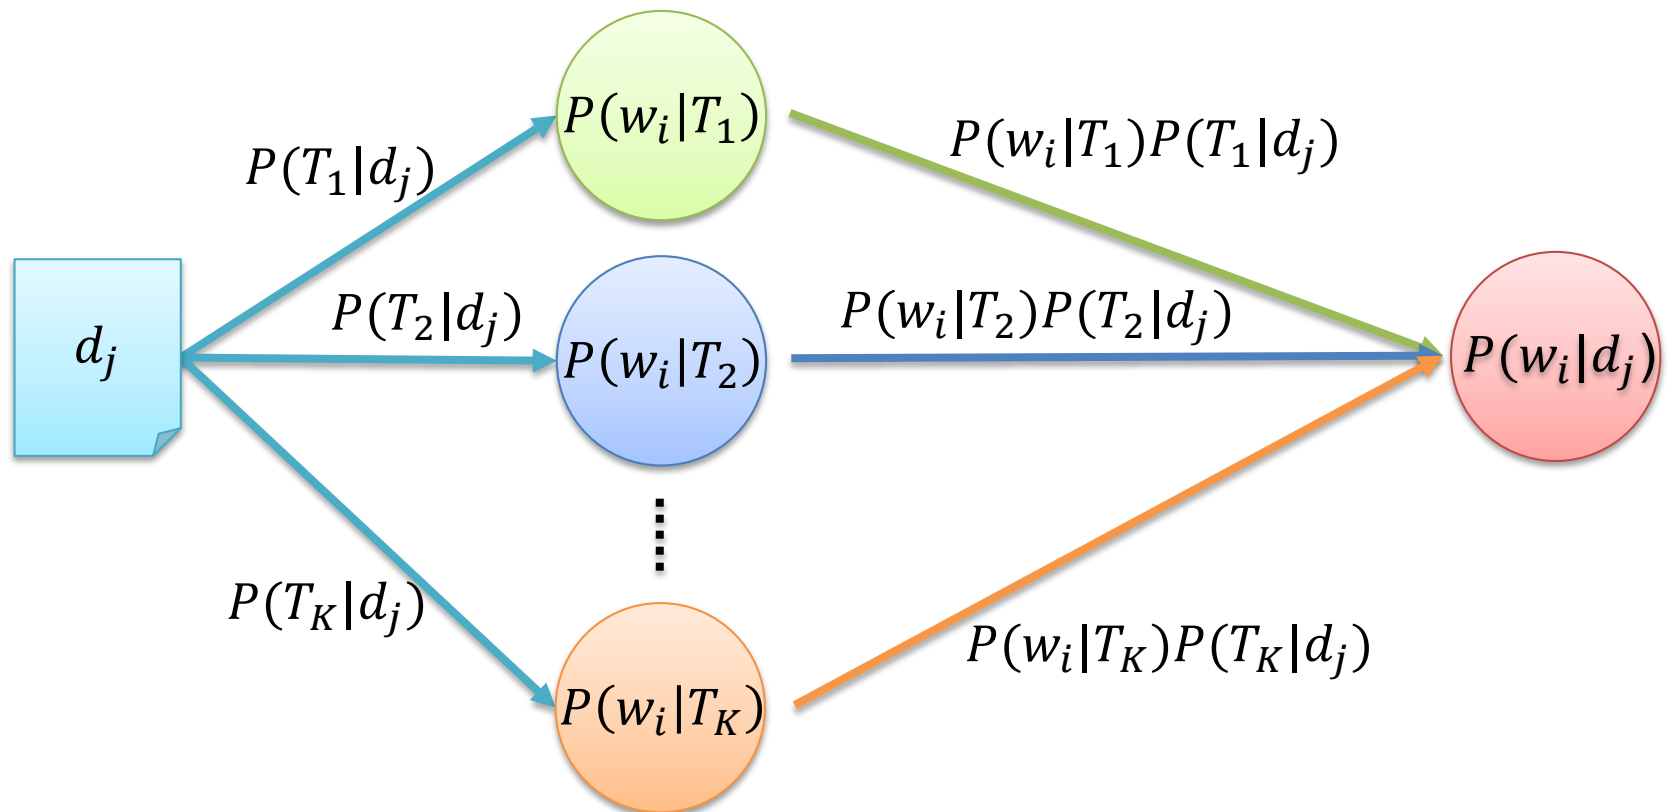
$$P(w_i, d_j) = P(d_j)P(w_i|d_j) = P(d_j) \sum_{k=1}^K P(w_i|T_k)P(T_k|d_j)$$

$$\begin{aligned} P(w_i|d_j) &= \sum_{k=1}^K P(w_i, T_k|d_j) = \sum_{k=1}^K \frac{P(w_i, T_k, d_j)}{P(d_j)} \\ &= \sum_{k=1}^K \frac{P(w_i, d_j|T_k)P(T_k)}{P(d_j)} \\ &= \sum_{k=1}^K \frac{P(w_i|T_k)P(d_j|T_k)P(T_k)}{P(d_j)} \\ &= \sum_{k=1}^K \frac{P(w_i|T_k)P(d_j, T_k)}{P(d_j)} = \sum_{k=1}^K P(w_i|T_k)P(T_k|d_j) \end{aligned}$$

Conditional Independence Assumption
document and word are independent
conditioned on the state of the associated
latent variable

PLSA – 2

- Thus, the modeling goal is to identify conditional probability mass functions $P(w_i|T_k)$ such that the document-specific word distributions $P(w_i|d_j)$ are as faithfully as possible approximated by convex combinations of these aspects



PLSA – 3

- The training objective is defined to maximize the total log-likelihood of a given training collection
 - The model parameters are $P(d_j)$, $P(w_i|T_k)$, and $P(T_k|d_j)$

$$\begin{aligned}\mathcal{L} &= \sum_{w_i \in V} \sum_{d_j \in D} c(w_i, d_j) \log P(w_i, d_j) \\ &= \sum_{w_i \in V} \sum_{d_j \in D} c(w_i, d_j) \log \left(P(d_j) \sum_{k=1}^K P(w_i|T_k) P(T_k|d_j) \right)\end{aligned}$$

PLSA – 4

- By using the Expectation-Maximization algorithm
 - E-step

$$P(T_k | w_i, d_j) = \frac{P(w_i | T_k) P(T_k | d_j)}{\sum_{k=1}^K P(w_i | T_k) P(T_k | d_j)}$$

- M-step

$$P(w_i | T_k) = \frac{\sum_{d_j \in \mathbf{D}} c(w_i, d_j) P(T_k | w_i, d_j)}{\sum_{i'=1}^{|V|} \sum_{d_j \in \mathbf{D}} c(w_{i'}, d_j) P(T_k | w_{i'}, d_j)}$$

$$P(T_k | d_j) = \frac{\sum_{i=1}^{|V|} c(w_i, d_j) P(T_k | w_i, d_j)}{\sum_{i'=1}^{|V|} c(w_{i'}, d_j)}$$

PLSA – 5

- Consequently, for a given word sequence, w_1, w_2, \dots, w_T , the joint probability in a language can be calculated by using PLSA

$$\begin{aligned} P(w_1, w_2, \dots, w_T) &= P(w_1) \prod_{t=2}^T P(w_t | w_1, w_2, \dots, w_{t-1}) \\ &= P(w_1) \prod_{t=2}^T \left(\sum_{k=1}^K P(w_t | T_k) P(T_k | w_1, w_2, \dots, w_{t-1}) \right) \end{aligned}$$

- Usually, we can combine the PLSA with the traditional n-gram models
 - Semantic matching and literal term matching

$$\begin{aligned} P(w_t | w_1, w_2, \dots, w_{t-1}) &= \alpha \cdot P(w_t | w_{t-n+1}, \dots, w_{t-1}) + \\ &\quad (1 - \alpha) \cdot \sum_{k=1}^K P(w_t | T_k) P(T_k | w_1, w_2, \dots, w_{t-1}) \end{aligned}$$

PLSA – 6

- For a new history of words, $w_1, w_2, \dots, w_{t-1} = H_1^{t-1}$, the **fold-in** strategy can be performed to obtain the topic distribution

- The word distribution for each topic $P(w_i|T_k)$ is fixed

- E-step

$$P(T_k|w_i, H_1^{t-1}) = \frac{P(w_i|T_k)P(T_k|H_1^{t-1})}{\sum_{k=1}^K P(w_i|T_k)P(T_k|H_1^{t-1})}$$

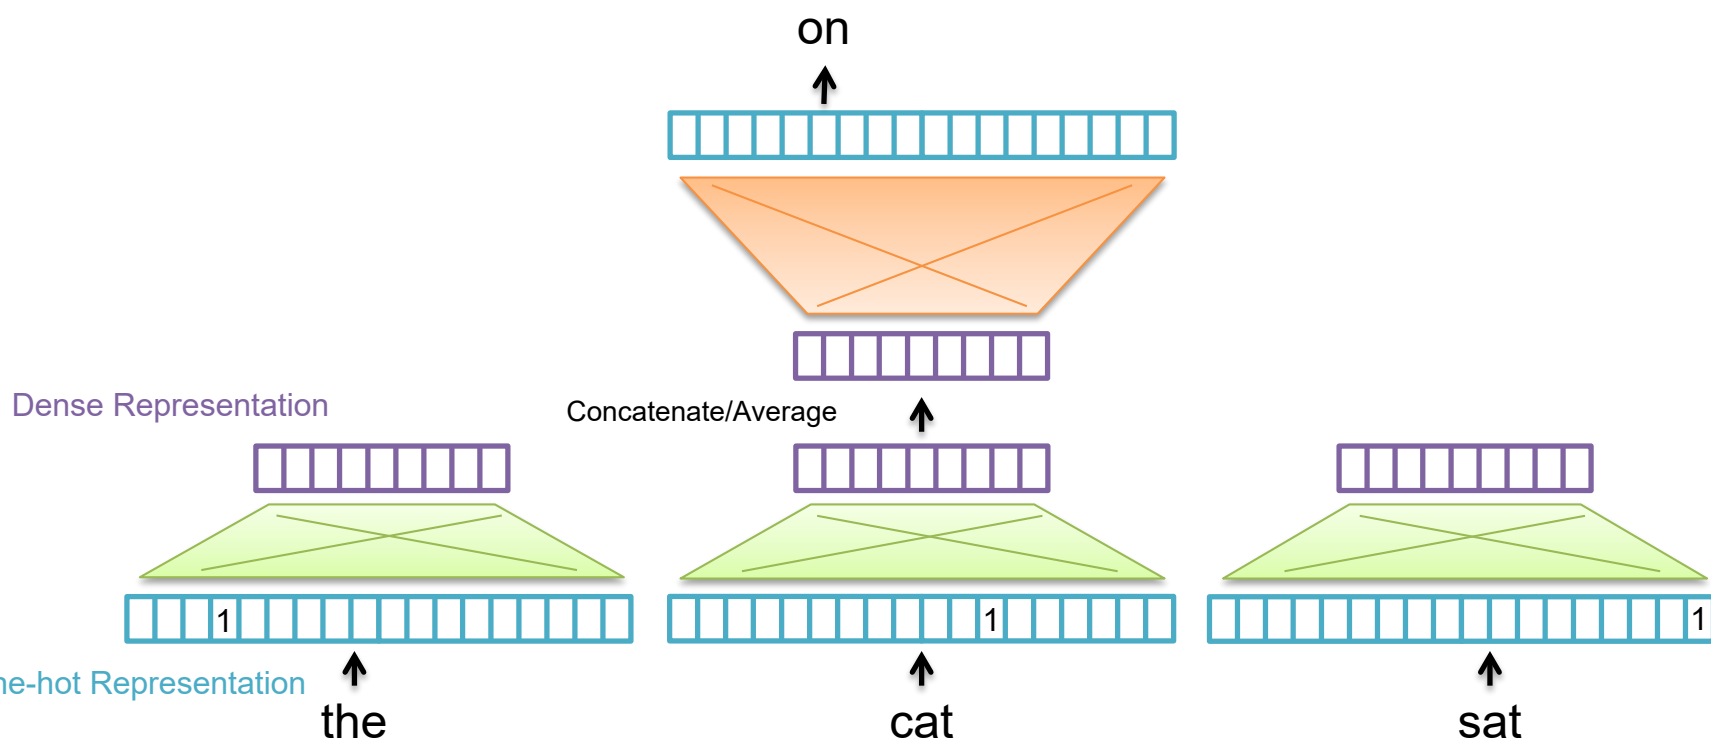
- M-step

$$P(T_k|H_1^{t-1}) = \frac{\sum_{i=1}^{|V|} c(w_i, H_1^{t-1})P(T_k|w_i, H_1^{t-1})}{\sum_{i'=1}^{|V|} c(w_{i'}, H_1^{t-1})}$$

Revisiting NNLM – 1

- The Neural Network Language Model (NNLM) estimated a statistical (n -gram) language model for **predicting future words**

$$P(w_1, w_2, \dots, w_T) \approx \prod_{t=1}^T P(w_t | w_{t-n+1}, \dots, w_{t-1})$$




Revisiting NNLM – 2.

$P(\text{there are books on the table})$

$\approx \mathbf{P(\text{there})}P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{are books})$

$P(\text{the}|\text{books on})P(\text{table}|\text{on the})$




Revisiting NNLM – 2..

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{are books})$

$P(\text{the}|\text{books on})P(\text{table}|\text{on the})$

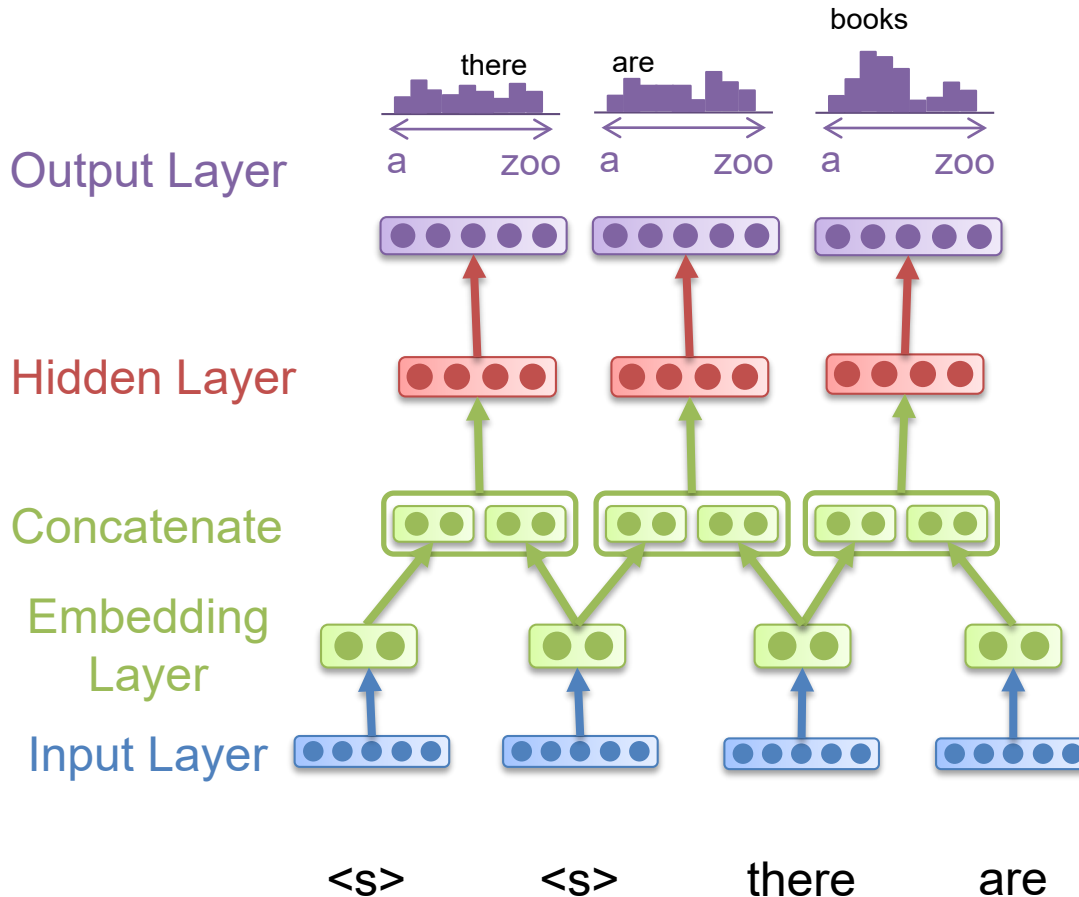


Revisiting NNLM – 2...

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})\mathbf{P(\text{books}|\text{there are})}P(\text{on}|\text{are books})$

$P(\text{the}|\text{books on})P(\text{table}|\text{on the})$

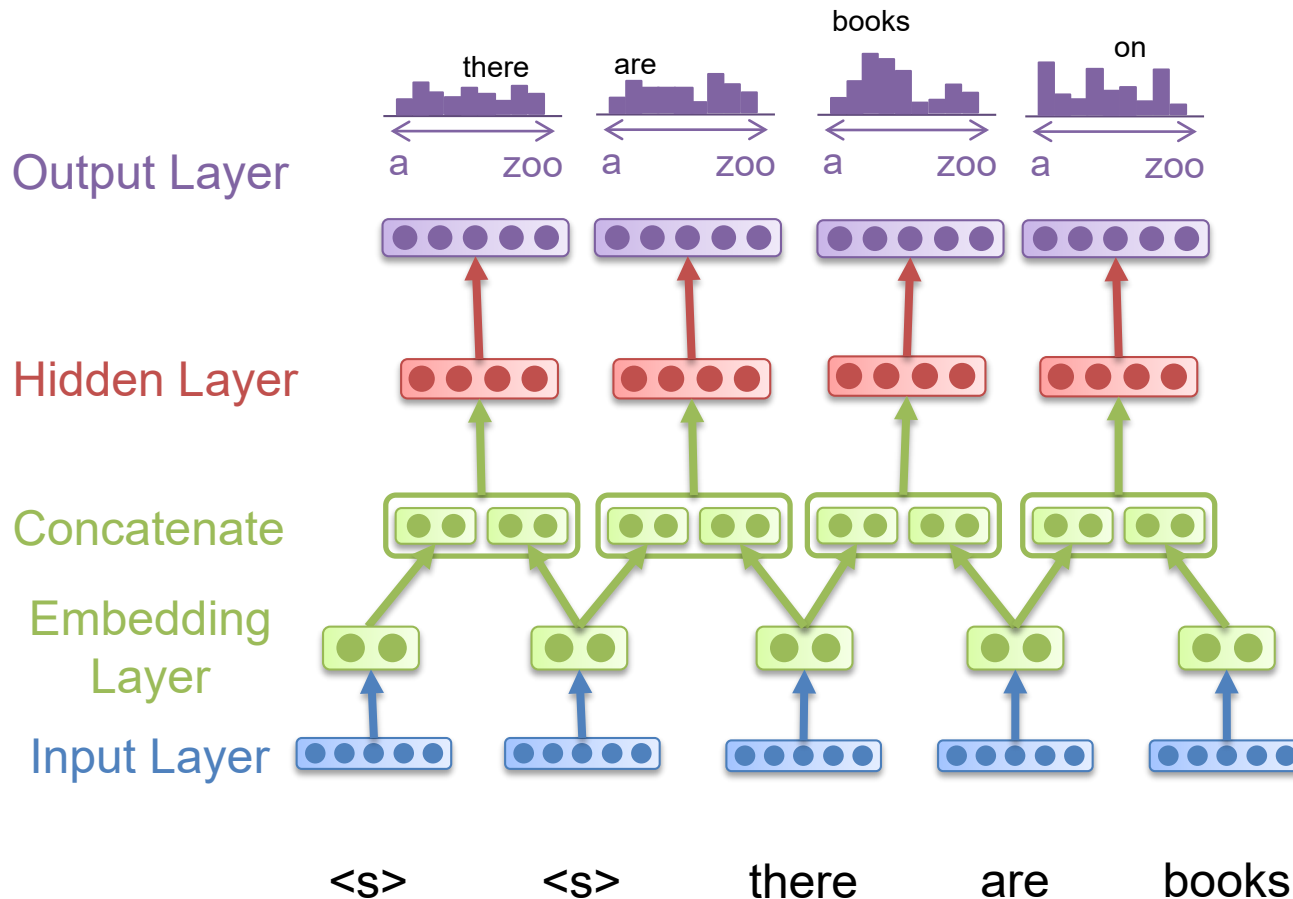


Revisiting NNLM – 2....

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})$ **$P(\text{on}|\text{are books})$**

$P(\text{the}|\text{books on})P(\text{table}|\text{on the})$

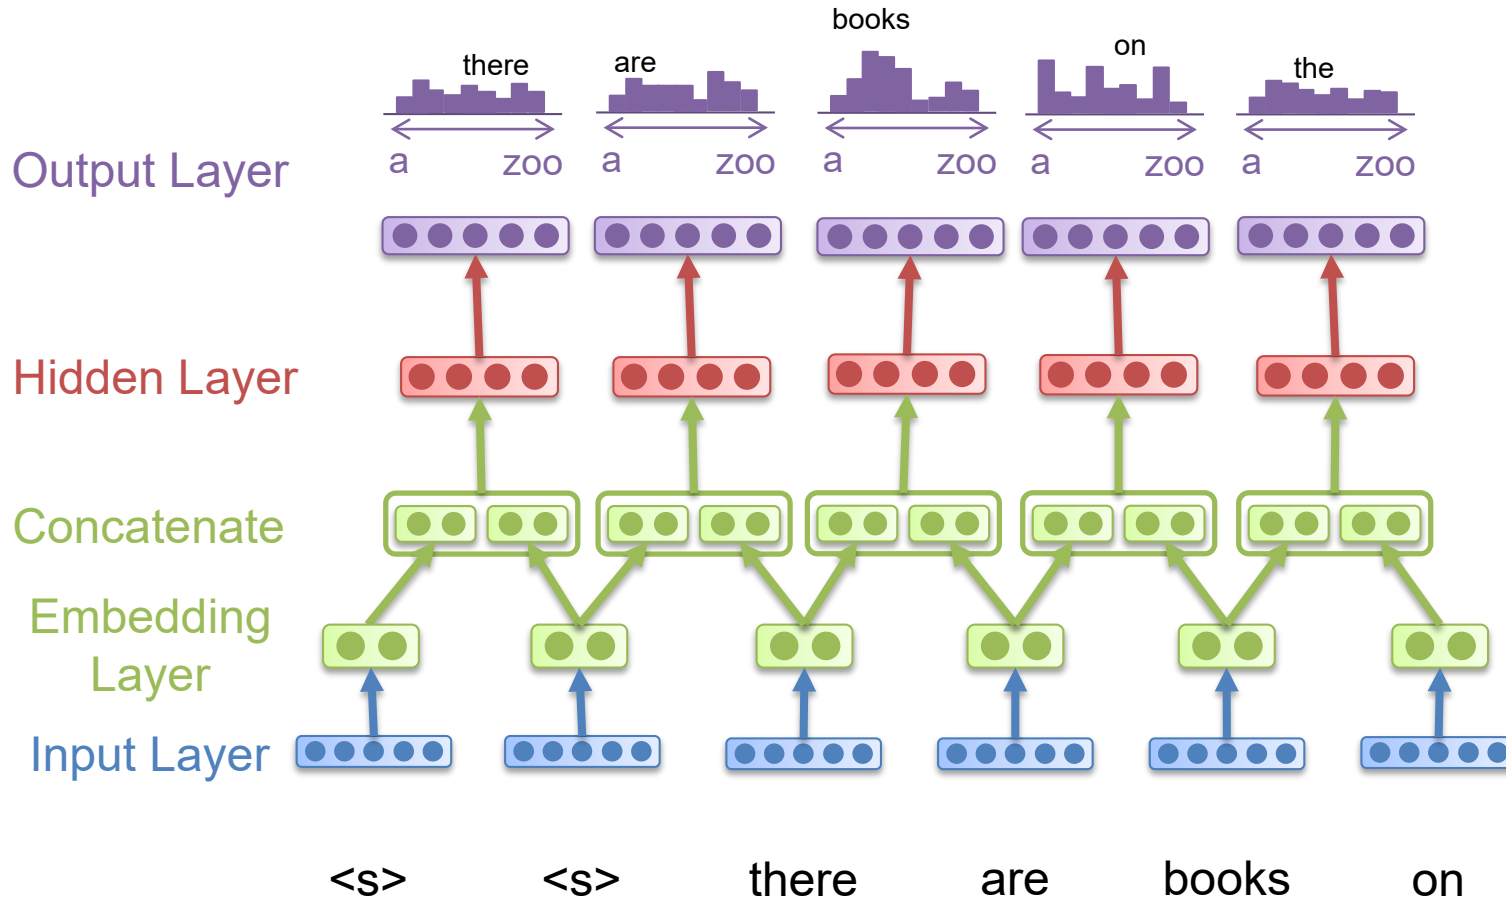


Revisiting NNLM – 2.....

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{are books})$

$P(\text{the}|\text{books on})P(\text{table}|\text{on the})$

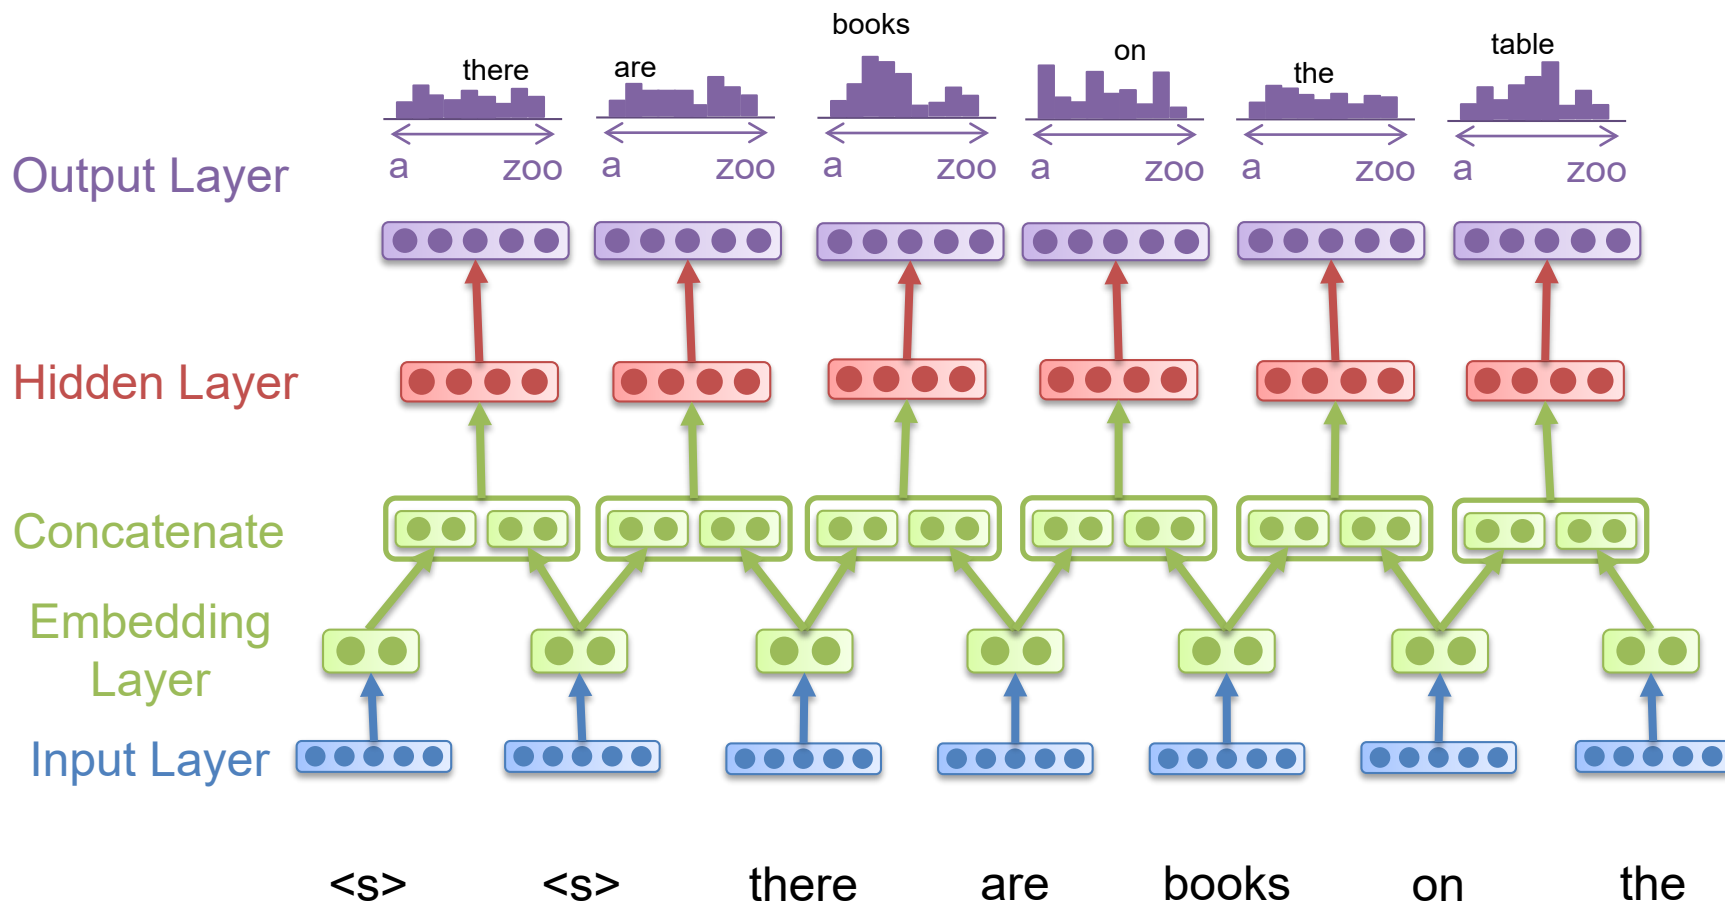


Revisiting NNLM – 2.....

$P(\text{there are books on the table})$

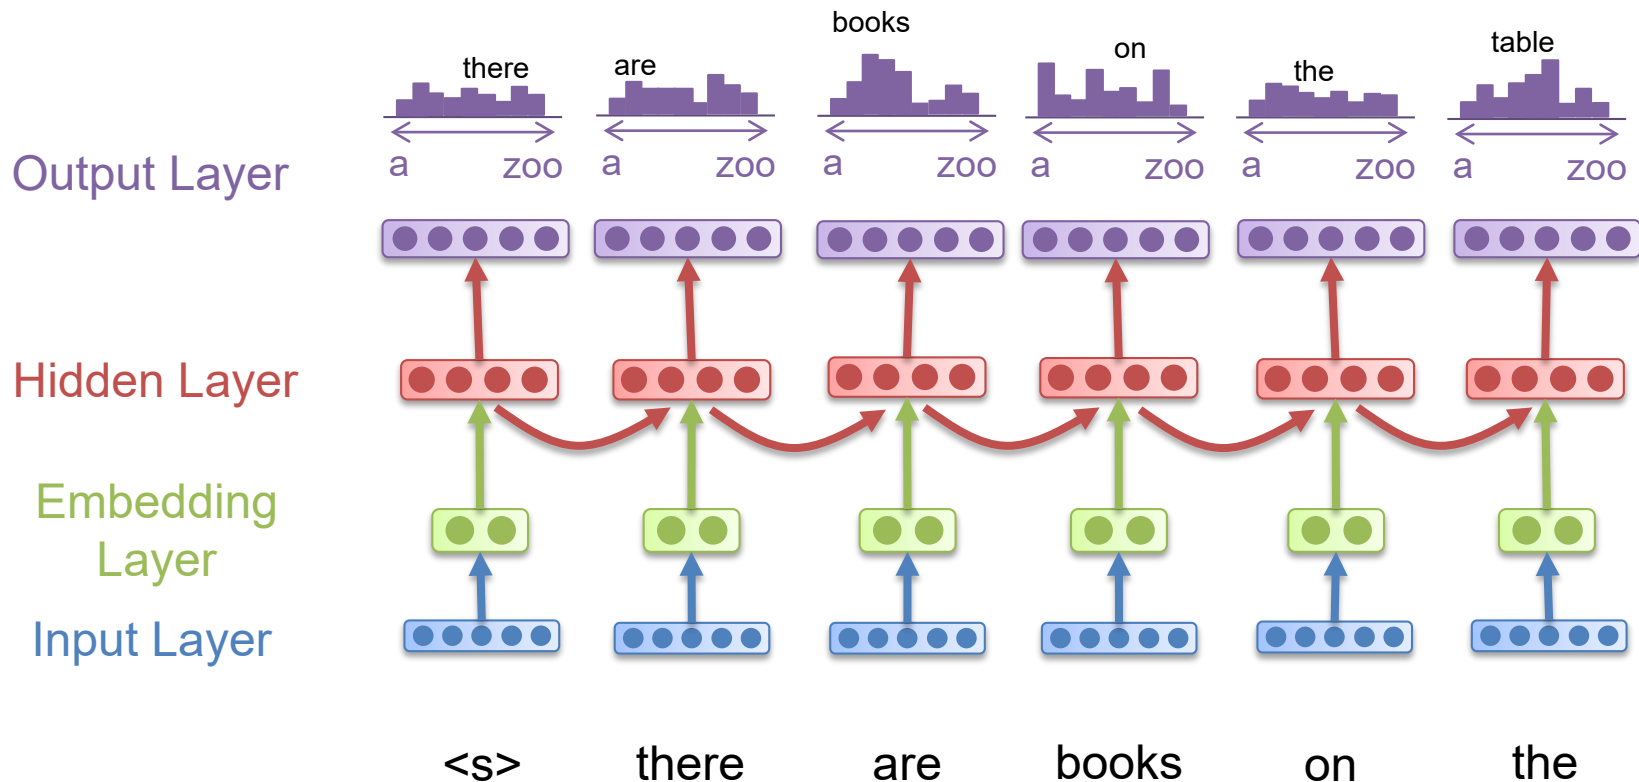
$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{are books})$

$P(\text{the}|\text{books on})\mathbf{P(\text{table}|\text{on the})}$



From NNLM to RNNLM

- The hidden state can encapsulate the information of word usage (ordering)
 - Leverage the information!!




RNNLM.

$P(\text{there are books on the table})$

$\approx \mathbf{P(\text{there})}P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{there are books})$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$




RNNLM..

$P(\text{there are books on the table})$

$\approx P(\text{there})\mathbf{P(\text{are}|\text{there})}P(\text{books}|\text{there are})P(\text{on}|\text{there are books})$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$




RNNLM...

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})\mathbf{P(\text{books}|\text{there are})}P(\text{on}|\text{there are books})$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$




RNNLM....

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})\mathbf{P(\text{on}|\text{there are books})}$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$




RNNLM.....

$P(\text{there are books on the table})$

$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{there are books})$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$

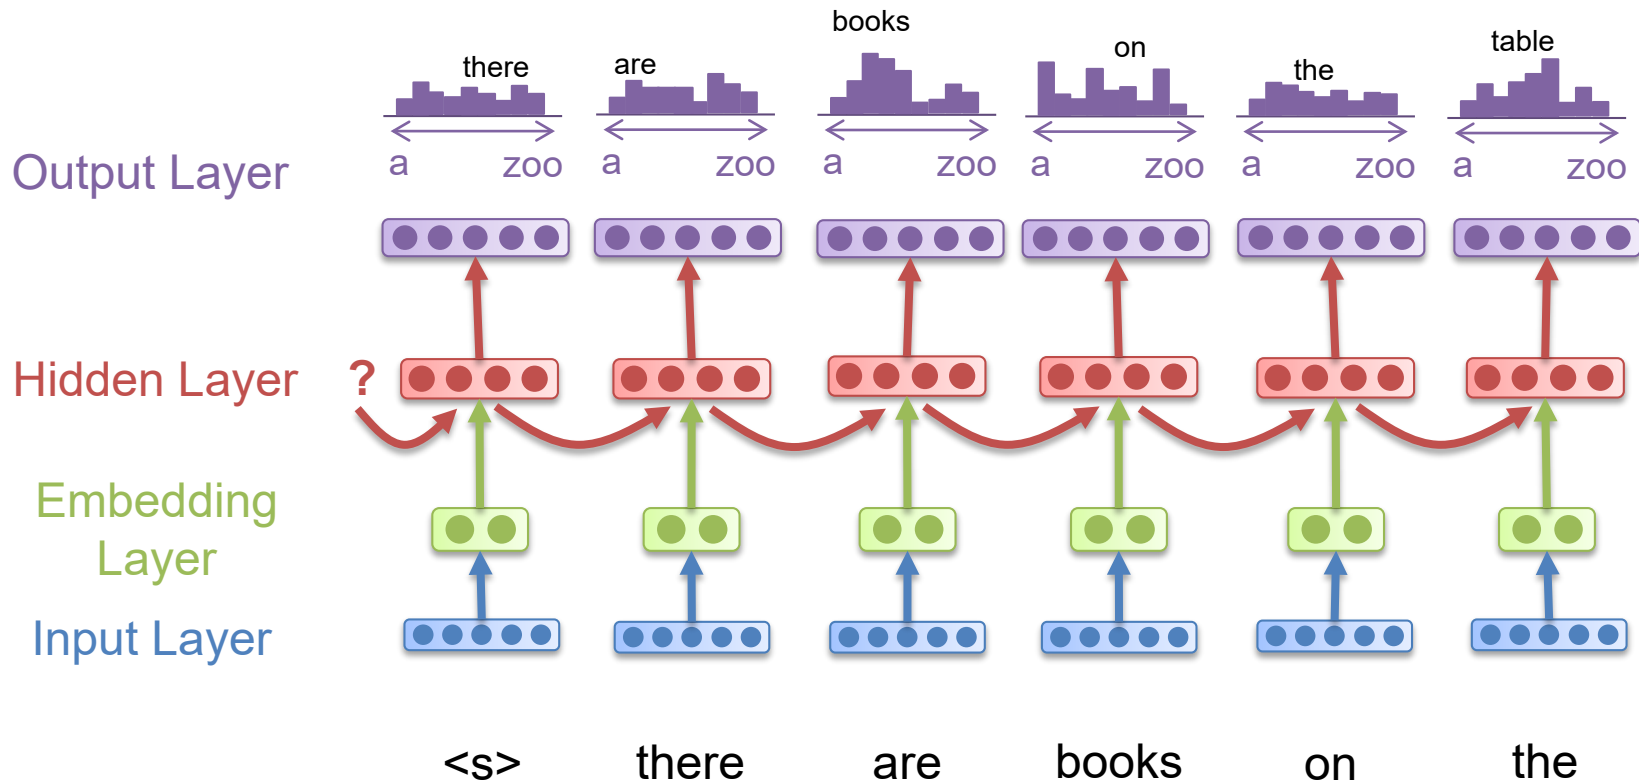


RNNLM.....

$P(\text{there are books on the table})$


$\approx P(\text{there})P(\text{are}|\text{there})P(\text{books}|\text{there are})P(\text{on}|\text{there are books})$

$P(\text{the}|\text{there are books on})P(\text{table}|\text{there are books on the})$



Recurrent Neural Network LM

- RNNLM has recently emerged as a promising modeling framework for several tasks
 - Both **word usage cues** and **long-span structural information** of word co-occurrence relationships can be taken into account naturally
- The limitations of the feed-forward NNLM
 - Need to specify the context length
 - RNN can efficiently represent more complex patterns than shallow NNs





Compared with Topic Modeling

	RNNLM	Topic Models
Local and/or Long-span Information	Both bigram and long-span information	Long-span
Capture the Long-span Information	By the Hidden State	By EM Algorithm
The Combination Weight between Local and Long-span Information	Automatic Adaptation	Empirical Setting
The Importance of Each Word in the History	Automatic Learned	Equal Weight
Interpretability	No	Yes

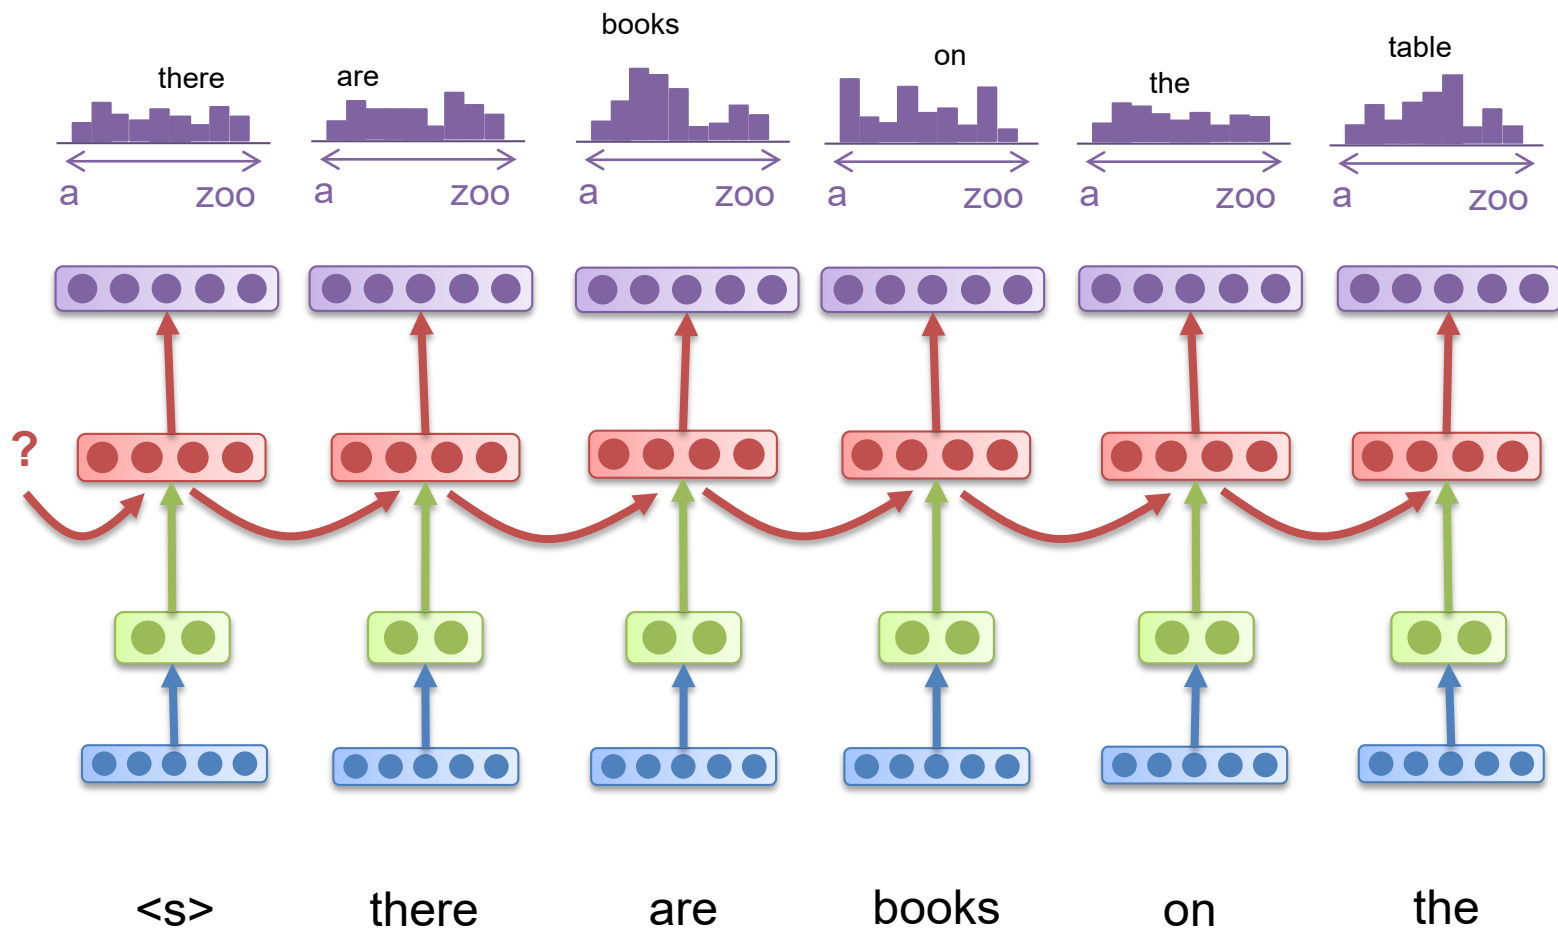
Elman & Jordan Types

Elman Type RNN
SimpleRNN
Vanilla RNN

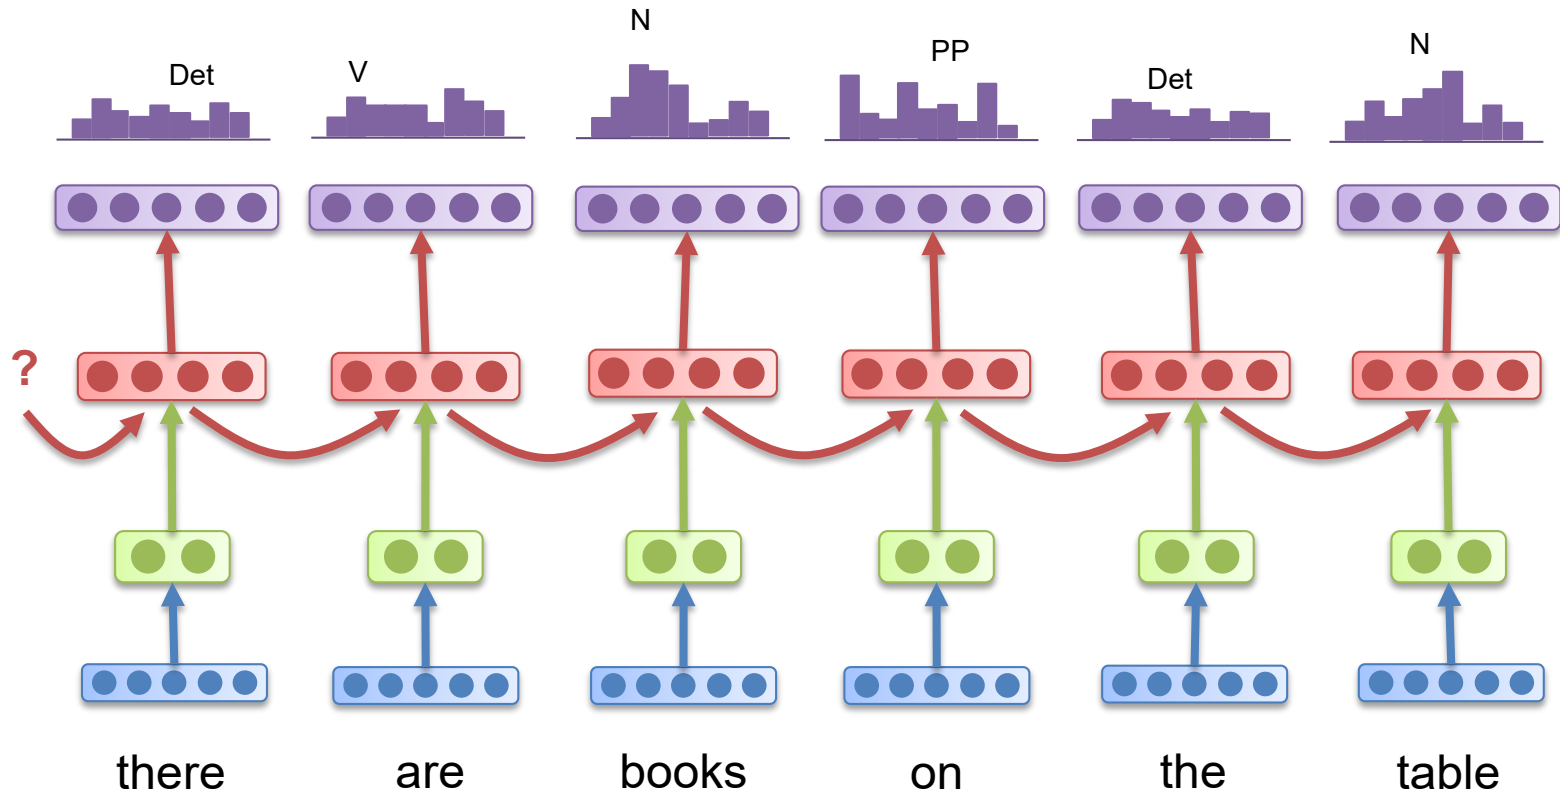
Jordan Type RNN



RNN for LM

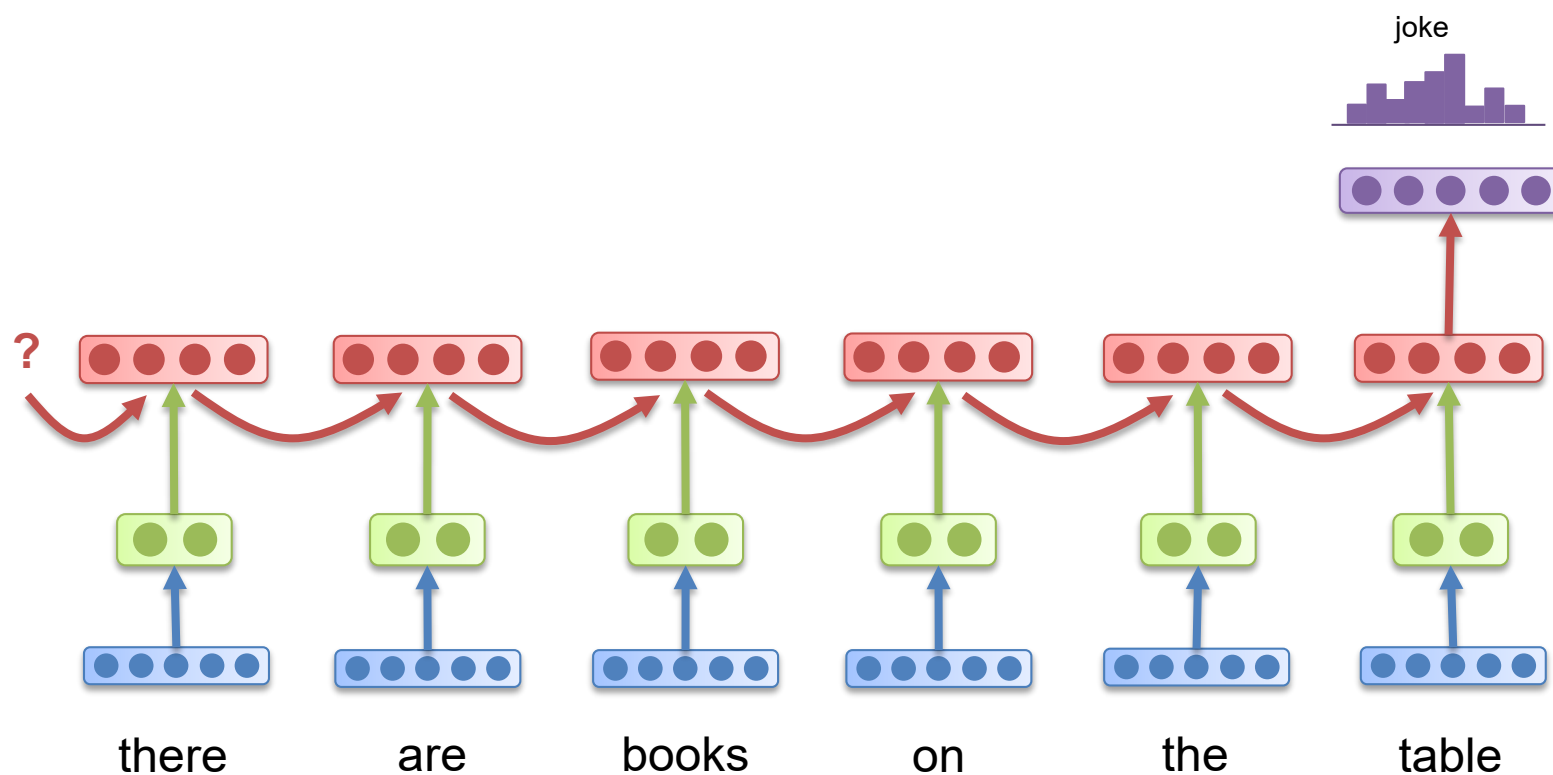


RNN for Tagging

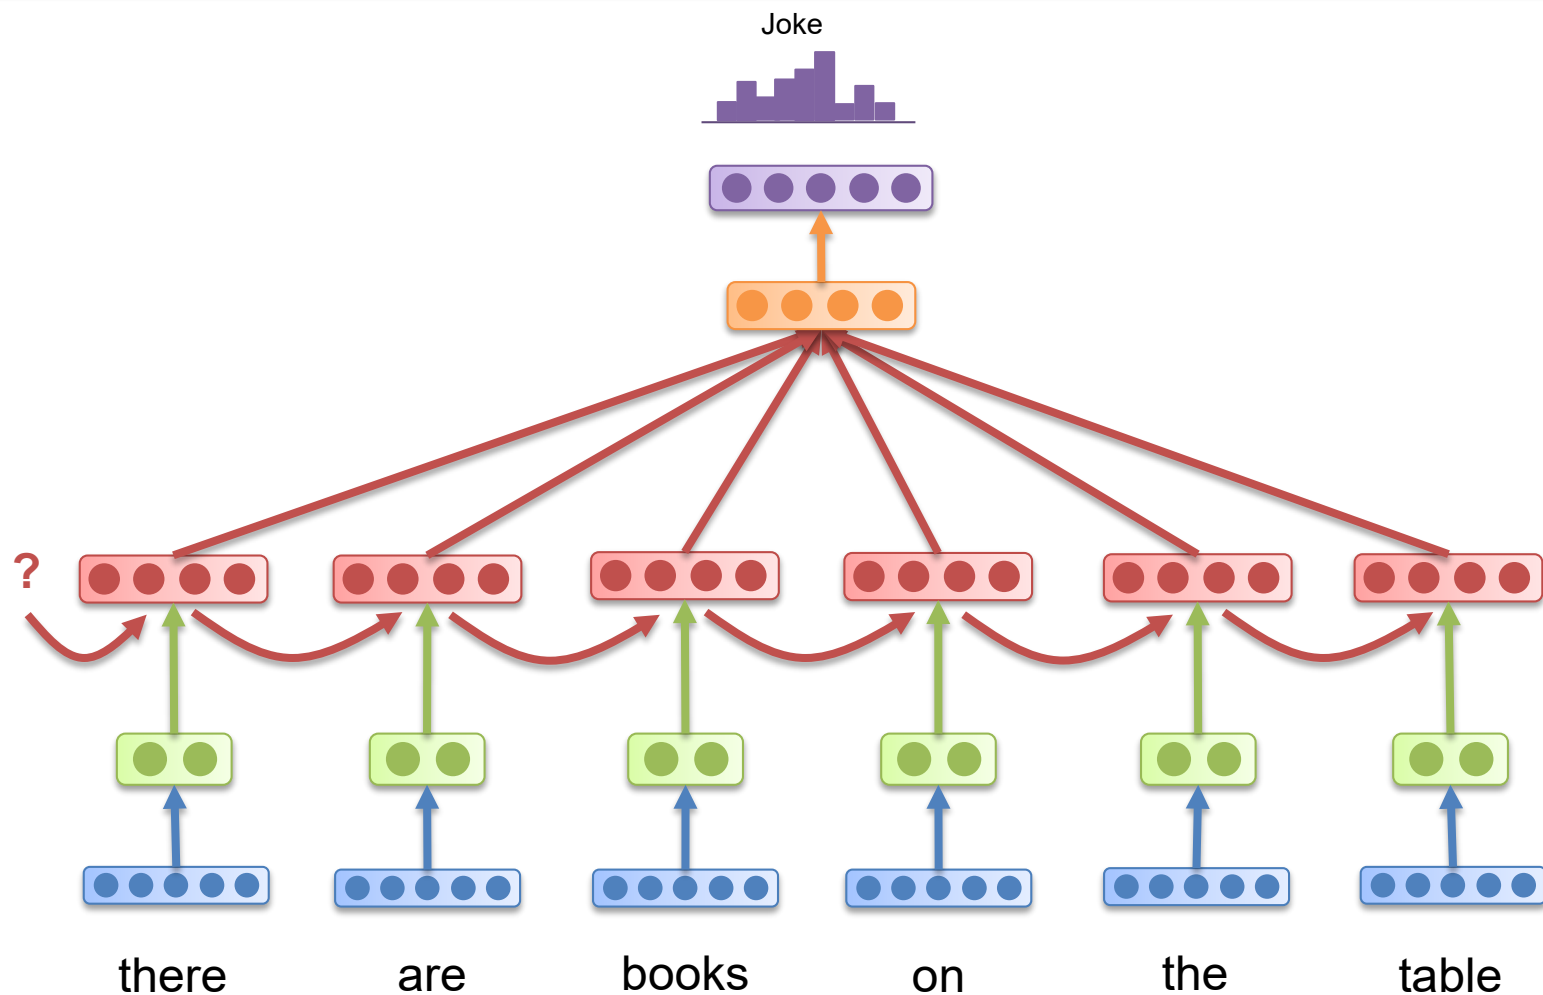


RNN for Classification – 1

- Forward RNN

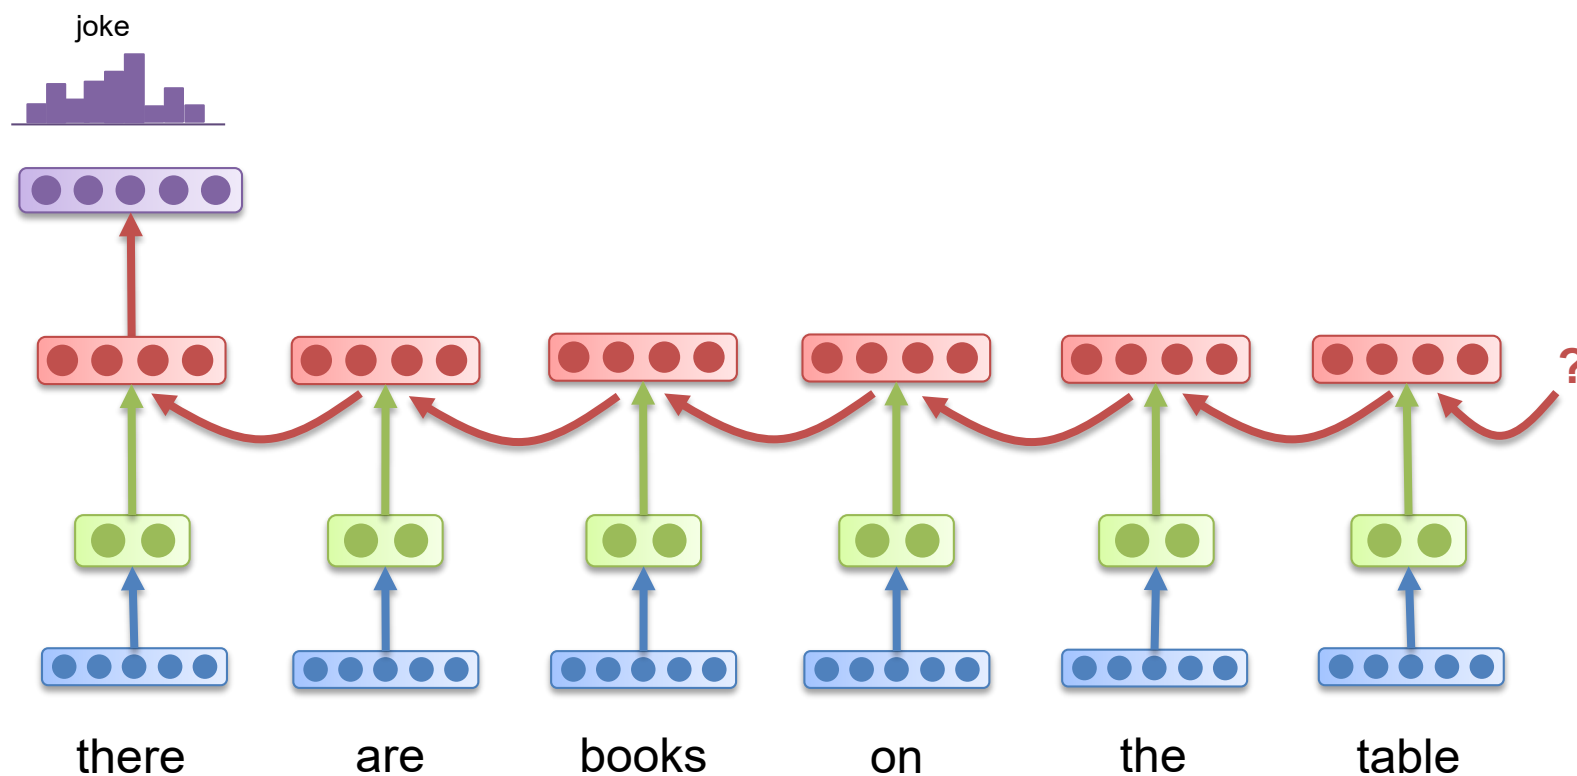


RNN for Classification – 2



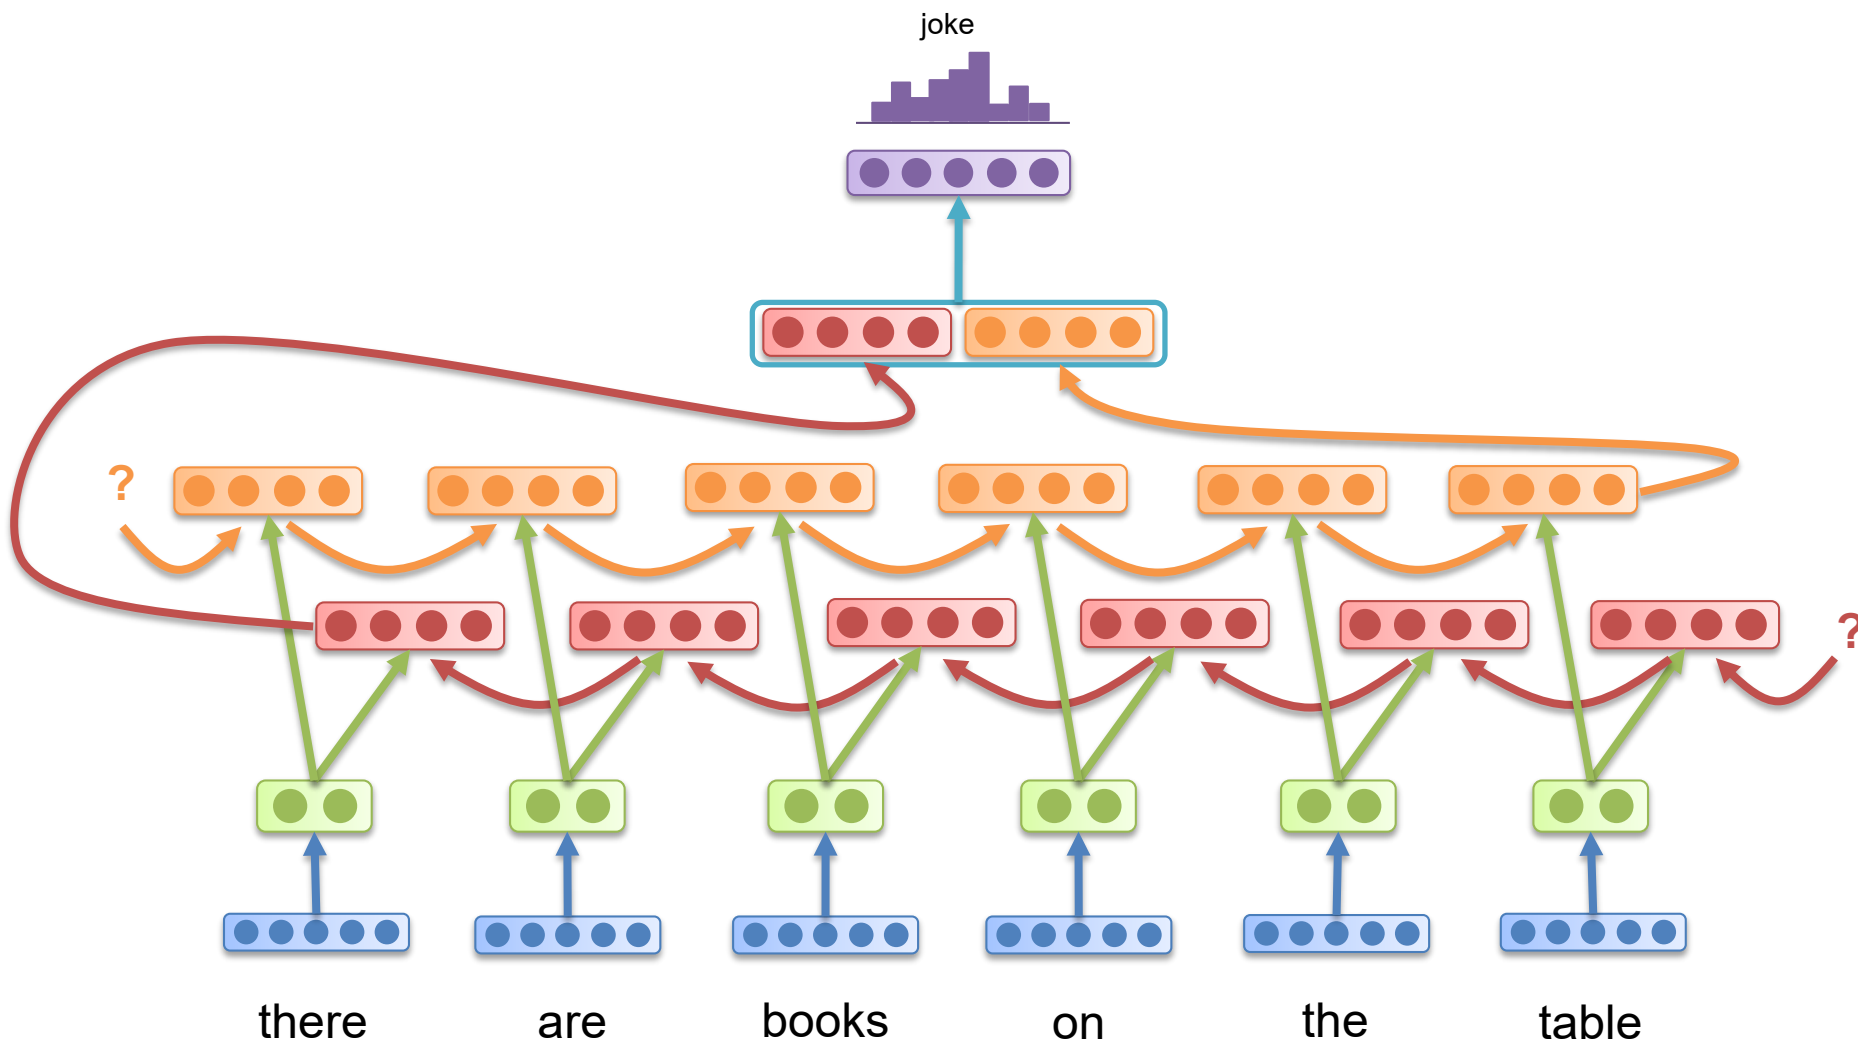
RNN for Classification – 3

- Backward RNN



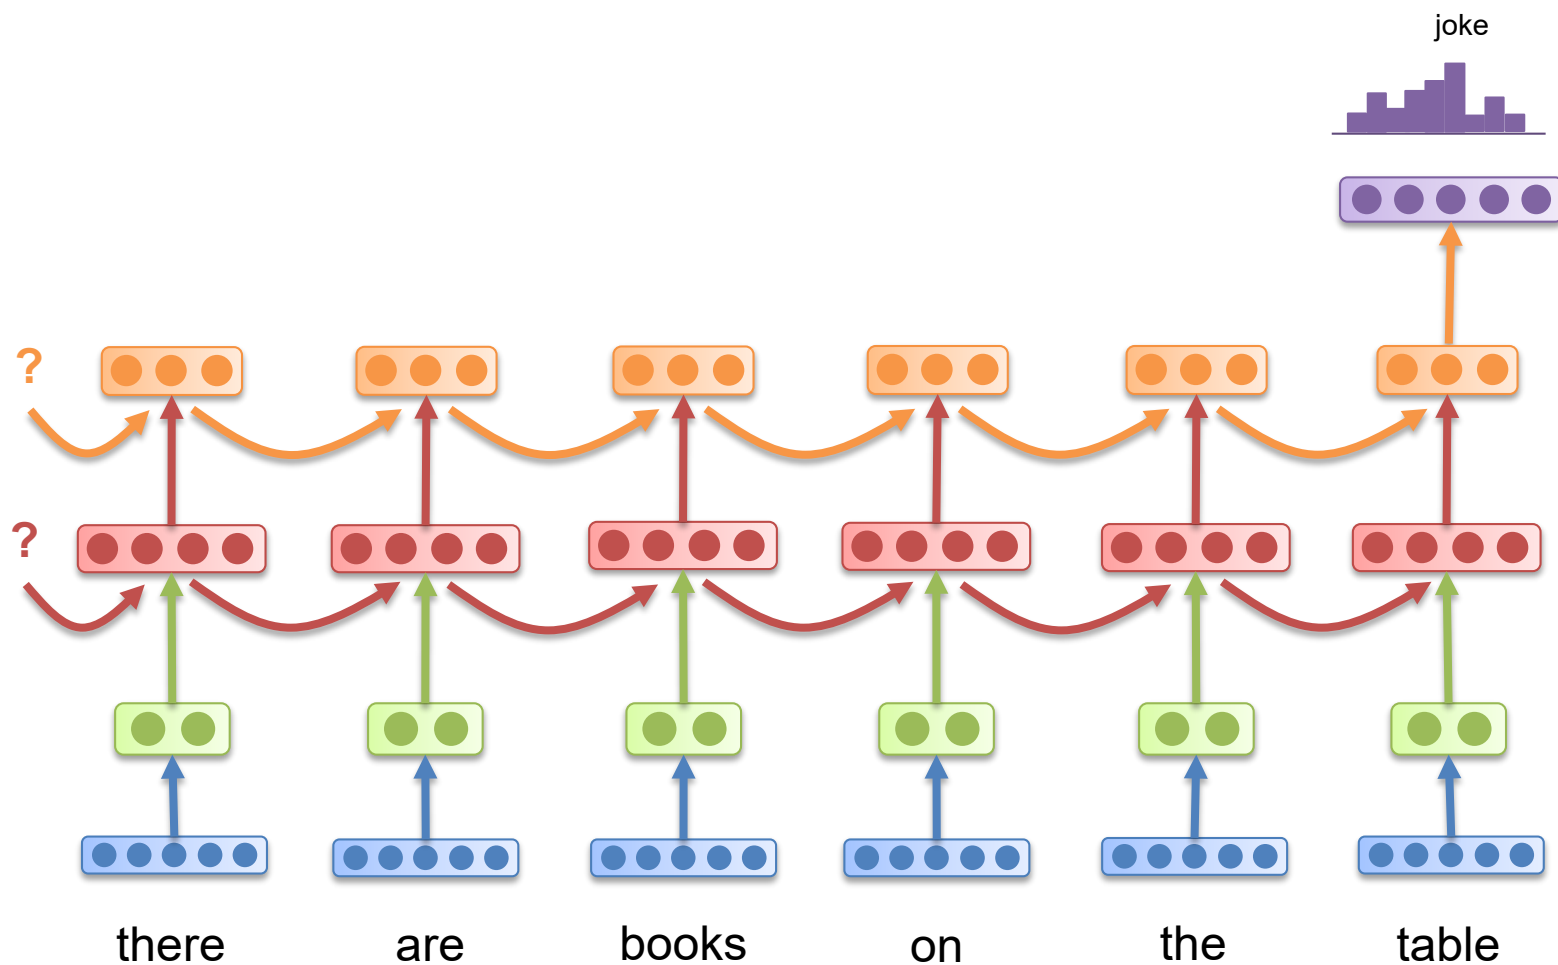
RNN for Classification – 4

- Bi-directional RNN!!




RNN for Classification – 5

- Multi-Layers RNN




Long Short-Term Memory (LSTM)

- Learning to Forget!
 - RNN
 - The classic model





- LSTM
 - Learning to forget
 - Capture longer information
 - Very slow in practice



Vanilla RNN

- RNN is hard to capture long-term dependencies

$$h_t = \tanh(W [h_{t-1}, x_t] + b)$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer





Concatenate



Copy

LSTM.

- The key to LSTMs is the **cell state**
 - The horizontal line running through the top of the diagram
 - It's very easy for information to just flow along it **unchanged**



Neural Network
Layer



Pointwise
Operation



Vector
Transfer



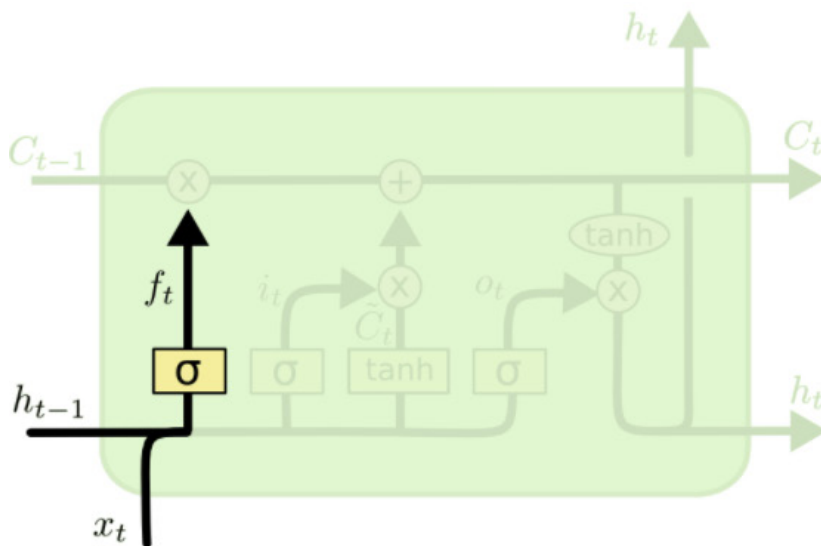
Concatenate




Copy

LSTM..

- The **forget gate** is to decide what information we're going to throw away from the cell state
 - $f_t = 1$: completely keep the information
 - $f_t = 0$: completely get rid of the information



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer




Concatenate




Copy

LSTM...

- The **input gate** is to decide which value we will update
 - A candidate vector, which contains the new information, will also be created



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer



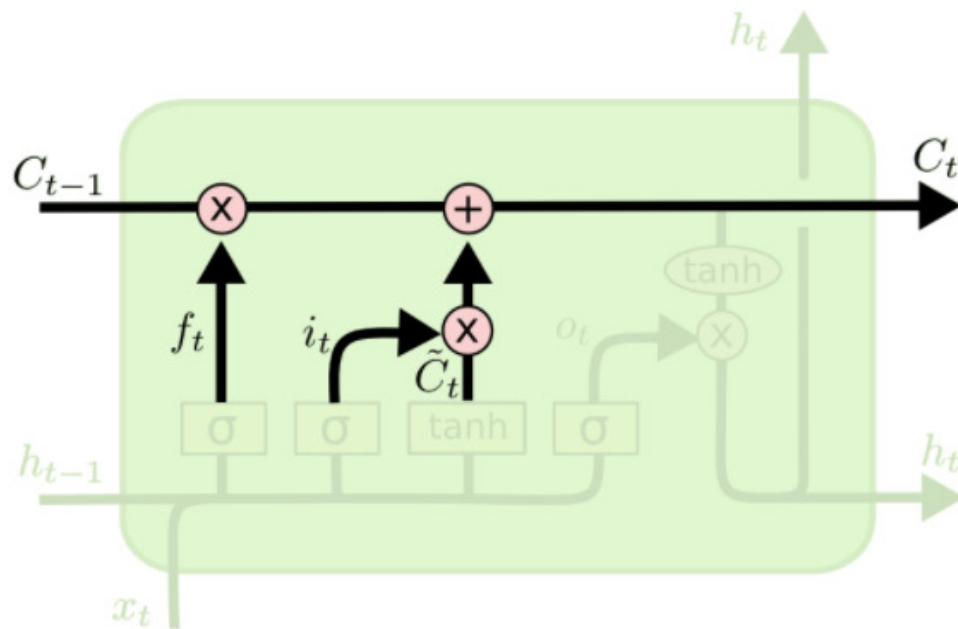
Concatenate




Copy

LSTM....

- Update the cell state!



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer




Concatenate



Copy


LSTM.....

- The **output gate** is a filter to select what information the model is going to output



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer






Concatenate



Copy


Gated Recurrent Unit (GRU)

- RNN
 - The classic model
- LSTM
 - Learning to forget
 - Capture longer information
 - Very slow in practice
- GRU
 - A balanced choice



GRU

- GRU combines the forget and input gates into a **update gate**
- It also merges the cell state and hidden state, and a **reset gate** is used to control the previous information




$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t]) \text{ Update Gate}$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t]) \text{ Reset Gate}$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$



Neural Network
Layer



Pointwise
Operation



Vector
Transfer



Concatenate



Copy

Questions?



kychen@mail.ntust.edu.tw