

# IMAGE-TO-IMAGE TRANSLATION WITH CONDITIONAL ADVERSARIAL NETWORKS



NATURE LANGUAGE PROCESSING

TA: 顏芷峯

DATE:5/17

# ABSTRACT

We investigate **conditional adversarial networks** as a general-purpose solution to image-to-image translation problems





# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion

# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion



# INTRODUCTION

- Many problems in **image processing, computer graphics, and computer vision** can be posed as “translating” an input image into a corresponding output image. Just as a concept may be expressed in either **English or French**

# INTRODUCTION

- we define **automatic image-to-image translation** as the task of translating one possible representation of a scene into another, given sufficient training data
- we explore **GANs in the conditional** setting. Just as GANs learn a generative model of data, conditional GANs learn a conditional generative model



# INTRODUCTION- GOAL

- Our goal in this paper is to **develop a common framework for all these problems.**

# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion




# CONTENT

- Introduction
- **Related work**
- Method
- Experiments
- Conclusion

# STRUCTURED LOSSES FOR IMAGE MODELING

- Image-to-image translation problems are often formulated as **per-pixel classification** or regression
- These formulations **treat the output space as “unstructured”** in the sense that each output pixel is considered conditionally independent from all others given the input image. **Conditional GANs instead learn a structured loss.** Structured losses penalize the joint configuration of the output.





## RELATED WORK

- Our method also differs from the prior works in several architectural choices for the generator and discriminator. Unlike past work, for our **generator** we use a “**U-Net**”-based architecture
- Our **discriminator** we **use** a convolutional “**PatchGAN**” classifier



# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion

# CONTENT

- Introduction
- Related work
- **Method**
- Experiments
- Conclusion



# OBJECTIVE

- **x**: input image, **z**: noise image, **y**: real image

# OBJECTIVE

- GAN  $G: z \rightarrow y$
- conditional GANs  $G: \{x, z\} \rightarrow y$



# OBJECTIVE

## Conditional GAN

$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y}[\log D(x, y)] + \mathbb{E}_{x,z}[\log(1 - D(x, G(x, z)))]$$

$$G^* = \operatorname{argmin}_G \max_D \mathcal{L}_{cGAN}(G, D)$$

# OBJECTIVE

$$\mathcal{L}_{GAN}(G, D) = \mathbb{E}_y[\log D(y)] + \mathbb{E}_{x,z}[\log(1 - D(x, G(x, z)))]$$

$$G^* = \operatorname{argmin}_G \max_D \mathcal{L}_{cGAN}(G, D)$$





•  $\mathcal{L}_{L_1}(G) = \mathbb{E}_{x,y,z}[\|y - G(x,z)\|_1]$

## OBJECTIVE\_FINAL OBJECTIVE

- $G^* = \operatorname{argmin}_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G)$



# NETWORK ARCHITECTURES

We adapt our generator and discriminator architectures  
Both generator and discriminator use modules of the form  
**convolution-BatchNorm-ReLu**

# Generator and Discriminator

The background is a gradient of dark blue and purple, speckled with white dots resembling a starry sky. On the right side, there are faint, light blue geometric patterns, including concentric circles and a circular scale with numerical markings from 0 to 210. In the bottom left corner, there are dashed circular lines with arrows indicating a clockwise direction.



# Generator and Discriminator



# TWO ARCHITECTURE CHOICES FOR THE GENERATOR

U-NET VS ENCODER



# TWO ARCHITECTURE CHOICES FOR THE GENERATOR


U-NET VS **ENCODER**

# ENCODER-DECODER

the input is passed through a series of layers

until a **bottleneck layer**, at which point the process is reversed

For many image translation problems, there is a great deal of low-level information shared between the input and output, and it would be desirable to shuttle this information directly across the net.






# TWO ARCHITECTURE CHOICES FOR THE GENERATOR

**U-NET** VS ENCODER


# U-NET

To give the generator a means to circumvent the bottleneck for information like this, we add skip connections, following the general shape of a “U-Net”. Specifically, we add skip connections between each layer  $i$  and layer  $n-i$  where  $n$  is the total number of layers. Each skip connection simply concatenates all channels at layer  $i$  with those at layer  $n-i$






# U-NET



From PAPER:U-Net: Convolutional Networks for Biomedical Image Segmentation

# U-NET



```
down1 = Conv2D(64, (3, 3), padding='same')(inputs)
```

```
down1 = BatchNormalization()(down1)
```

```
down1 = Activation('relu')(down1)
```

```
down1 = Conv2D(64, (3, 3), padding='same')(down1)
```

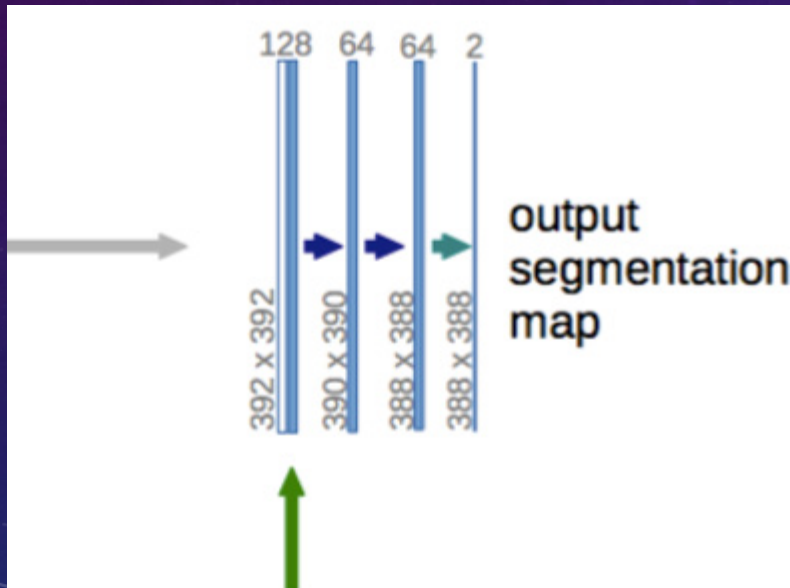
```
down1 = BatchNormalization()(down1)
```

```
down1 = Activation('relu')(down1)
```

```
down1_pool = MaxPooling2D((2, 2), strides=(2, 2))(down1)
```



# U-NET



```
up1 = UpSampling2D((2, 2))(up2)
```



```
up1 = concatenate([down1, up1], axis=3)
```

```
up1 = Conv2D(64, (3, 3), padding='same')(up1)
```

```
up1 = BatchNormalization()(up1)
```

```
up1 = Activation('relu')(up1)
```

```
up1 = Conv2D(64, (3, 3), padding='same')(up1)
```

```
up1 = BatchNormalization()(up1)
```

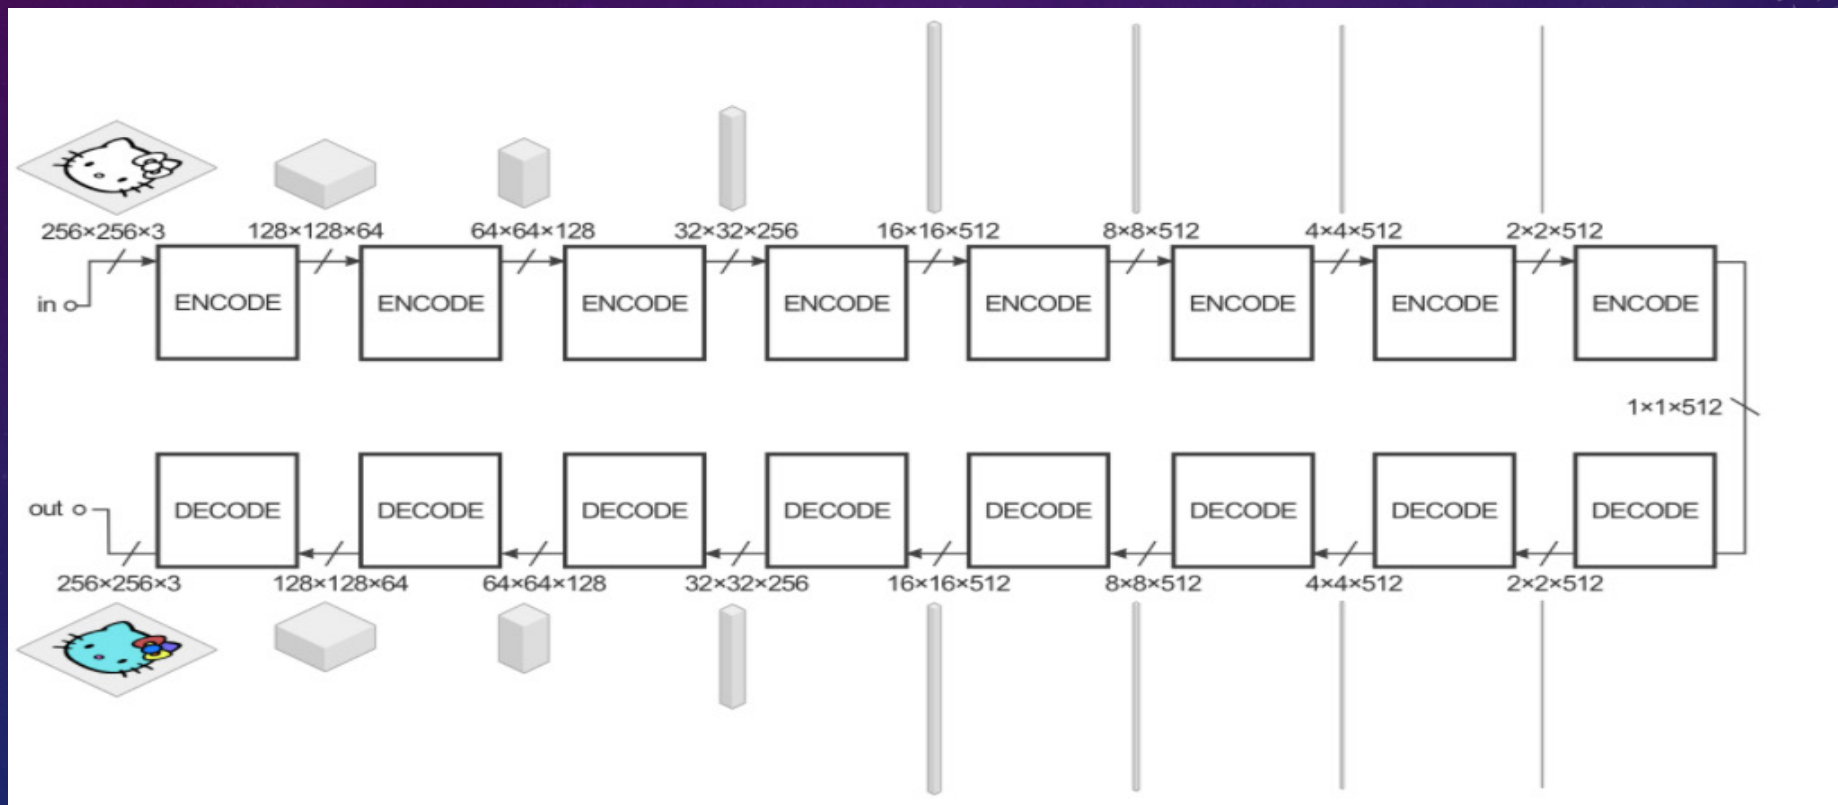
```
up1 = Activation('relu')(up1)
```

```
up1 = Conv2D(64, (3, 3), padding='same')(up1)
```

```
up1 = BatchNormalization()(up1)
```

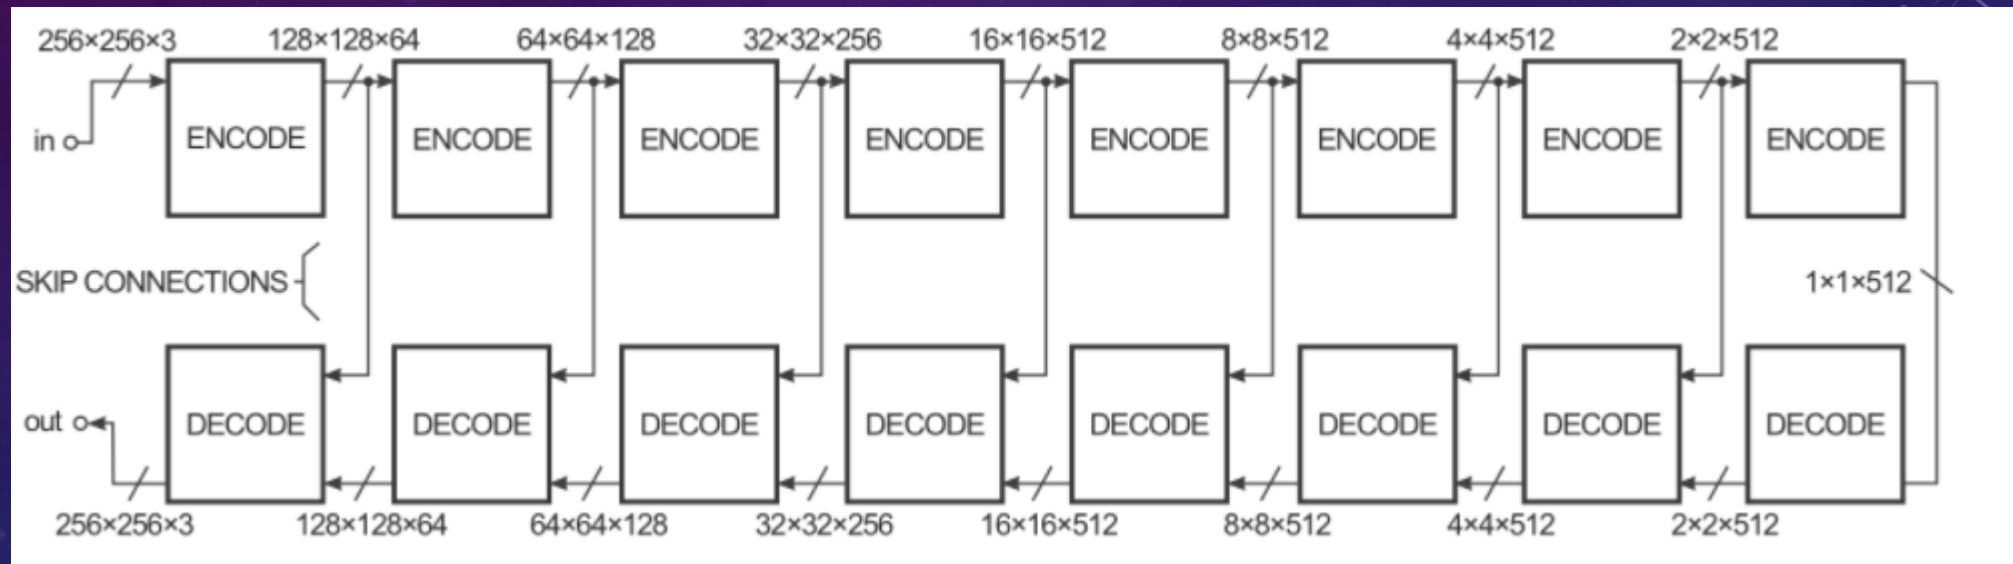
```
up1 = Activation('relu')(up1)
```

# ENCODER-DECODER






# U-NET




# Generator and Discriminator






# Generator and Discriminator



# MARCOV RANDOM FIELD

Result

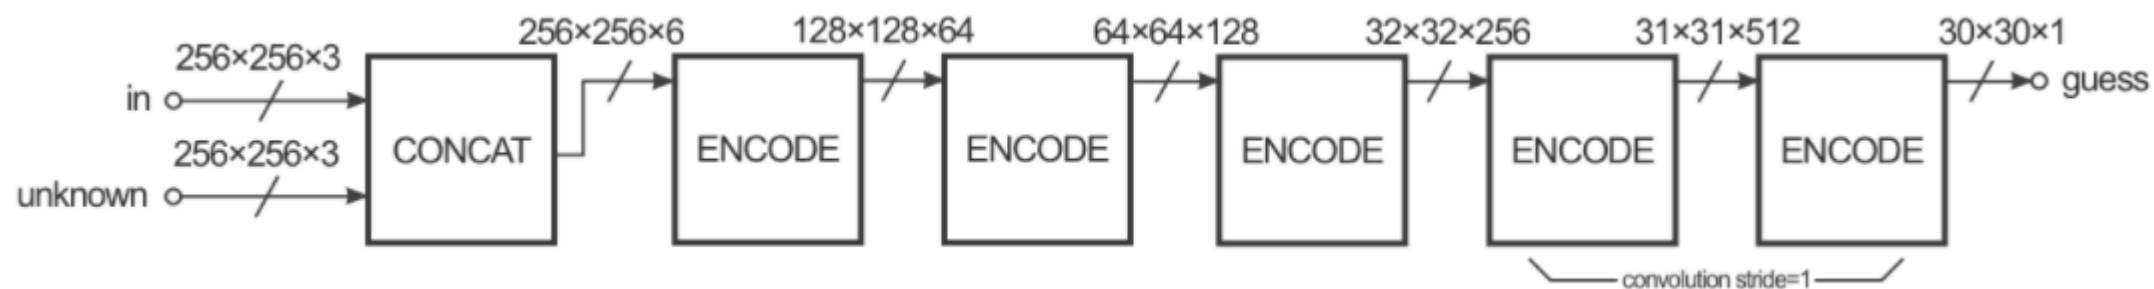


Input





# MARKOVIAN DISCRIMINATOR



# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion



# CONTENT

- Introduction
- Related work
- Method
- **Experiments**
- Conclusion

# EXPERIMENTS

- Semantic labels->photo, trained on the Cityscapes dataset
- Architectural labels->photo, trained on CMP Facades
- Map->aerial photo, trained on data scraped from Google Maps.
- BW->color photos
- Edges->photo using the HED edge detector plus postprocessing.
- Sketch->photo: tests edges!photo models on hand-drawn
- Day->night
- Photo with missing pixels->in-painted photo, trained on Paris StreetView



# EXPERIMENTS

Labels to Street Scene




input



output

Aerial to Map




input



output

Labels to Facade




input



output

BW to Color



input



output

Day to Night



input



output

Edges to Photo



input



output



# EXPERIMENTS



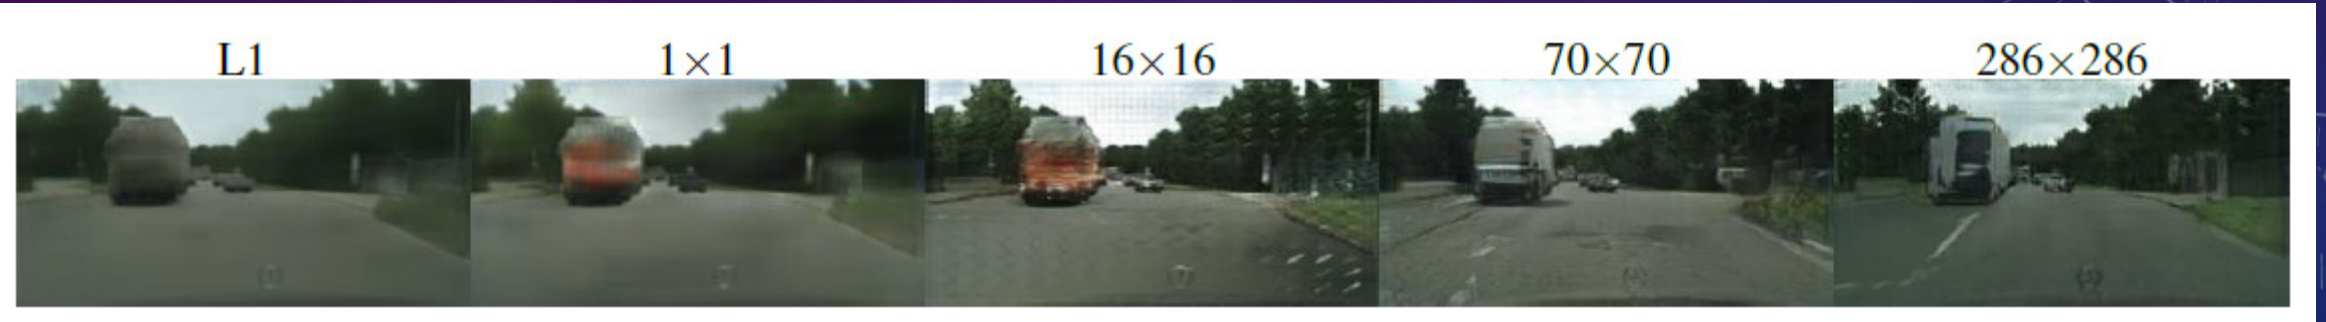


# EXPERIMENTS

Adding skip connections to an encoder-decoder to create a “U-Net” results in much higher quality results.



# EXPERIMENTS



Patch size variations. Uncertainty in the output manifests itself differently for different loss functions



# EXPERIMENTS

Loss	Per-pixel acc.	Per-class acc.	Class IOU
L1	0.42	0.15	0.11
GAN	0.22	0.05	0.01
cGAN	0.57	0.22	0.16
L1+GAN	0.64	0.20	0.15
L1+cGAN	<b>0.66</b>	<b>0.23</b>	<b>0.17</b>
Ground truth	0.80	0.26	0.21

Table 1: FCN-scores for different losses, evaluated on Cityscapes labels $\leftrightarrow$ photos.

Loss	Per-pixel acc.	Per-class acc.	Class IOU
Encoder-decoder (L1)	0.35	0.12	0.08
Encoder-decoder (L1+cGAN)	0.29	0.09	0.05
U-net (L1)	0.48	0.18	0.13
U-net (L1+cGAN)	<b>0.55</b>	<b>0.20</b>	<b>0.14</b>

Table 2: FCN-scores for different generator architectures (and objectives), evaluated on Cityscapes labels $\leftrightarrow$ photos. (U-net (L1+cGAN) scores differ from those reported in other tables since batch size was 10 for this experiment and 1 for other tables, and random variation between training runs.)

# EXPERIMENTS

Discriminator receptive field	Per-pixel acc.	Per-class acc.	Class IOU
<b>1×1</b>	0.39	0.15	0.10
<b>16×16</b>	0.65	0.21	<b>0.17</b>
<b>70×70</b>	<b>0.66</b>	<b>0.23</b>	<b>0.17</b>
<b>286×286</b>	0.42	0.16	0.11

Table 3: FCN-scores for different receptive field sizes of the discriminator, evaluated on Cityscapes labels→photos. Note that input images are  $256 \times 256$  pixels and larger receptive fields are padded with zeros.

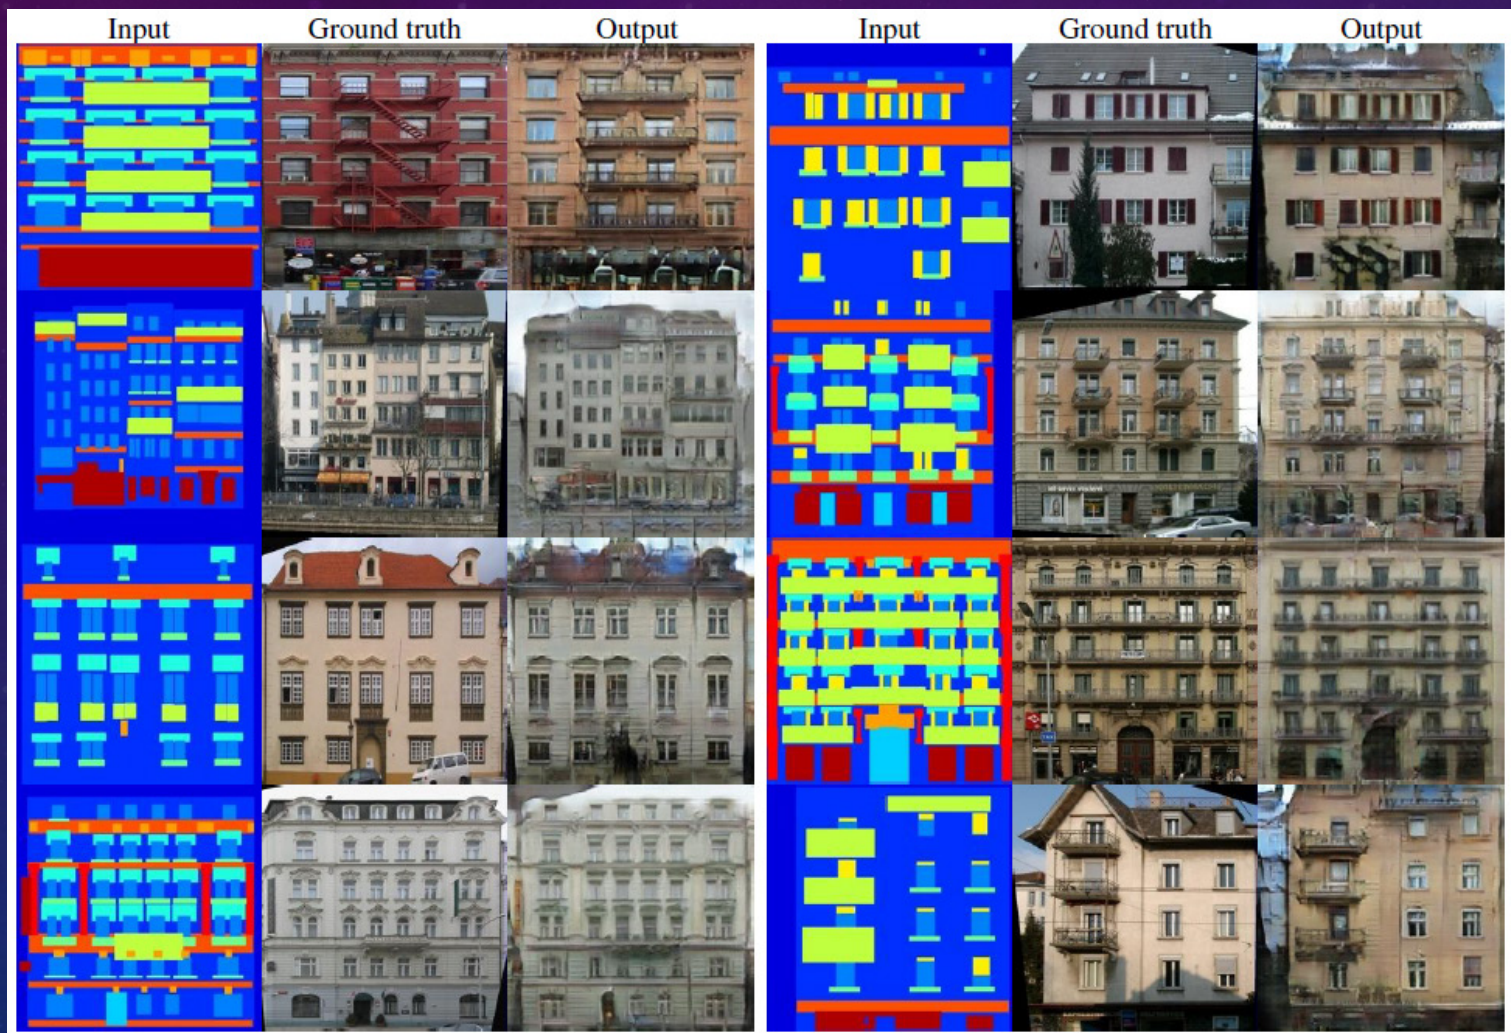


# EXPERIMENTS





# EXPERIMENTS






# EXPERIMENTS





# EXPERIMENTS





# EXPERIMENTS



# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion



# CONTENT

- Introduction
- Related work
- Method
- Experiments
- Conclusion

# RESULTS

- The results in this paper suggest that **conditional adversarial networks** are a promising approach **for** many **image-to-image translation** tasks, especially those involving highly structured graphical outputs. These networks learn a loss adapted to the task and data at hand, which makes them applicable in a wide variety of settings.